



Computer Viewing

CS 432 Interactive Computer
Graphics

Prof. David E. Breen

Department of Computer Science



Objectives

- Introduce the mathematics of projection
- Introduce OpenGL viewing functions
- Look at alternate viewing APIs



Computer Viewing

- There are three aspects of the viewing process, all of which should be / are implemented in the pipeline,
 - Positioning the camera
 - Setting the model-view matrix
 - Selecting a lens
 - Setting the projection matrix
 - Clipping
 - Setting the view volume




The OpenGL Camera

- In OpenGL, initially the object and camera frames are the same
 - Default model-view matrix is an identity
- The camera is located at origin and points in the negative z direction
- OpenGL also specifies a default view volume that is a cube with sides of length 2 centered at the origin
 - Default projection matrix is an identity



Default Projection

Default projection is orthographic





Moving the Camera Frame


- If we want to visualize object with both positive and negative z values we can either
 - Move the camera in the positive z direction
 - Translate the camera frame
 - Move the objects in the negative z direction
 - Translate the world frame
- Both of these views are equivalent and are determined by the model-view matrix
 - Want a translation (`Translate(0.0, 0.0, -d);`)
 - $d > 0$

Moving Camera back from Origin


frames after translation by $-d$

$$d > 0$$

default frames




(a)



(b)

Moving the Camera

- We can move the camera to any desired position by a sequence of rotations and translations
- Example: side view
 - Rotate the camera
 - Move it away from origin
 - Model-view matrix $C = TR$




OpenGL code

- Remember that last transformation specified is first to be applied

```
// Using mat.h
```

```
mat4 t = Translate(0.0, 0.0, -d);  
mat4 ry = RotateY(-90.0);  
mat4 m = t*ry;
```





The LookAt Function

- The GLU library contained the function `gluLookAt` to form the required modelview matrix through a simple interface
- Note the need for setting an up direction
 - Should not be parallel to look-at direction
- Replaced by `LookAt()` in `mat.h`
 - Can concatenate with modeling transformations
- Example: isometric view of cube aligned with axes

```
mat4 mv = LookAt(vec4 eye, vec4 at, vec4 up);
```


LookAt Function

`LookAt(eye, at, up)`

`eye` - location of camera

`at` - look-at point

`up` - up vector





Other Viewing APIs

- The LookAt function is only one possible API for positioning the camera
- Others include
 - View reference point, view plane normal, view up (PHIGS, GKS-3D)
 - Yaw, pitch, roll
 - Elevation, azimuth, twist
 - Direction angles



Projections and Normalization

- The default projection in the eye (camera) frame is orthographic
- For points within the default view volume

$$x_p = x$$

$$y_p = y$$

$$z_p = 0$$

- Most graphics systems use *view normalization*
 - All other views are converted to the default view by transformations that determine the projection matrix
 - Allows use of the same pipeline for all views



Homogeneous Coordinate Representation

default orthographic projection

$$\begin{aligned}x_p &= x \\y_p &= y \\z_p &= 0 \\w_p &= 1\end{aligned}$$


$$\mathbf{p}_p = \mathbf{M}\mathbf{p}$$

$$\mathbf{M} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

In practice, we can let $\mathbf{M} = \mathbf{I}$ and set the z term to zero later


Simple Perspective

- Center of projection at the origin
- Projection plane $z = d$, $d < 0$



Perspective Equations

Consider top and side views



$$x_p = \frac{x}{z/d}$$

$$y_p = \frac{y}{z/d}$$

$$z_p = d$$



Homogeneous Coordinate Form

consider $\mathbf{q} = \mathbf{M}\mathbf{p}$ where $\mathbf{M} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1/d & 0 \end{bmatrix}$

$$\mathbf{q} = \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} \Rightarrow \mathbf{p} = \begin{bmatrix} x \\ y \\ z \\ z/d \end{bmatrix}$$



Perspective Division

- However $w \neq 1$, so we must divide by w to return from homogeneous coordinates
- This *perspective division* yields


$$x_p = \frac{x}{z/d} \quad y_p = \frac{y}{z/d} \quad z_p = d$$

the desired perspective equations

- We will consider the corresponding clipping volume with mat.h functions that are equivalent to deprecated OpenGL functions

OpenGL Orthogonal Viewing


Ortho(left, right, bottom, top, near, far)



near and far measured from camera


OpenGL Perspective

`Frustum(left, right, bottom, top, near, far)`



Using Field of View

- With **Frustum** it is often difficult to get the desired view
- **Perspective (fovy, aspect, near, far)** often provides a better interface





Projection Matrices

CS 432 Interactive Computer Graphics

Prof. David E. Breen

Department of Computer Science



Objectives

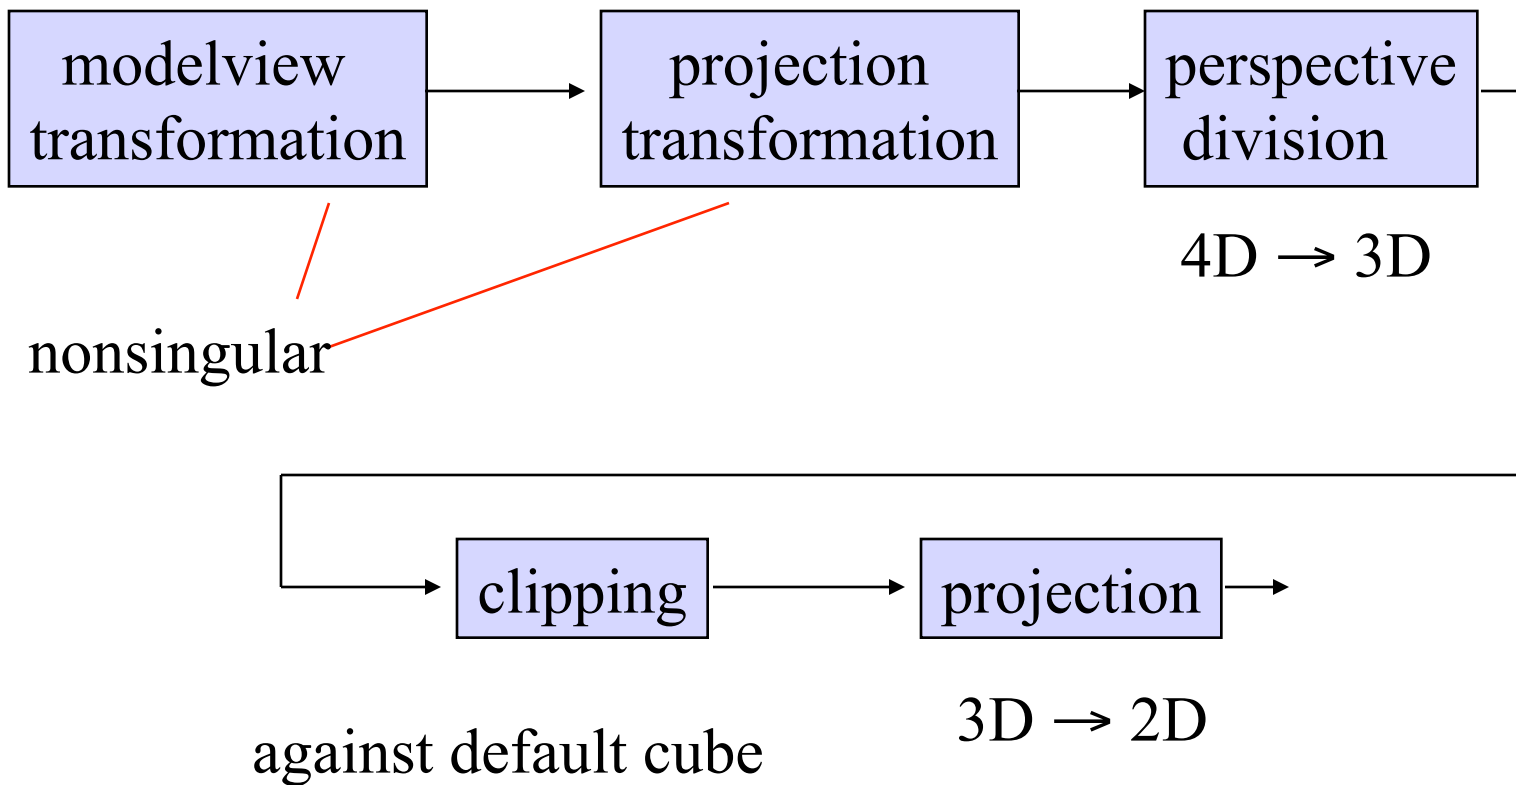
- Derive the projection matrices used for standard OpenGL projections
- Introduce oblique projections
- Introduce projection normalization



Normalization

- Rather than derive a different projection matrix for each type of projection, we can convert all projections to orthogonal projections with the default view volume
- This strategy allows us to use standard transformations in the pipeline and makes for efficient clipping

Pipeline View






Notes

- We stay in four-dimensional homogeneous coordinates through both the modelview and projection transformations
 - Both these transformations are nonsingular
 - Default to identity matrices (orthogonal view)
- Normalization lets us clip against simple cube regardless of type of projection
- Delay final projection until end
 - Important for hidden-surface removal to retain depth information as long as possible

Orthographic Normalization

`Ortho(left, right, bottom, top, near, far)`

normalization \Rightarrow find transformation to convert
specified clipping volume to default



$\text{left} < \text{right}$ $\text{bottom} < \text{top}$ $\text{near} < \text{far}$



Orthographic Matrix

- Two steps

- Move center to origin

$$T(-(left+right)/2, -(bottom+top)/2, (near+far)/2))$$

- Scale to have sides of length 2

$$S(2/(left-right), 2/(top-bottom), 2/(near-far))$$

$$\mathbf{P} = \mathbf{ST} = \begin{bmatrix} \frac{2}{right - left} & 0 & 0 & -\frac{right + left}{right - left} \\ 0 & \frac{2}{top - bottom} & 0 & -\frac{top + bottom}{top - bottom} \\ 0 & 0 & \frac{2}{near - far} & \frac{far + near}{far - near} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



Final Projection

- Set $z = 0$
- Equivalent to the homogeneous coordinate transformation


$$\mathbf{M}_{\text{orth}} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

- Hence, general orthographic projection in 4D is

$$\mathbf{P} = \mathbf{M}_{\text{orth}} \mathbf{S} \mathbf{T}$$


Oblique Projections

- The OpenGL projection functions cannot produce general parallel projections such as



- However if we look at the example of the cube it appears that the cube has been sheared
- Oblique Projection = Shear + Orthographic Projection

General Shear





Shear Matrix

xy shear (z values unchanged)


$$\mathbf{H}(\theta, \phi) = \begin{bmatrix} 1 & 0 & -\cot \theta & 0 \\ 0 & 1 & -\cot \phi & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Projection matrix

$$\mathbf{P} = \mathbf{M}_{\text{orth}} \mathbf{H}(\theta, \phi)$$


General case: $\mathbf{P} = \mathbf{M}_{\text{orth}} \mathbf{STH}(\theta, \phi)$

Equivalency



Effect on Clipping


- The projection matrix $\mathbf{P} = \mathbf{STH}$ transforms the original clipping volume to the default clipping volume



Simple Perspective

Consider a simple perspective with the COP at the origin, the near clipping plane at $z = -1$, and a 90 degree field of view determined by the planes

$$x = \pm z, y = \pm z$$





Perspective Matrices

Simple projection matrix in homogeneous coordinates

$$\mathbf{M} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

Note that this matrix is independent of the far clipping plane

Generalization

$$\mathbf{N} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & \alpha & \beta \\ 0 & 0 & -1 & 0 \end{bmatrix}$$

after perspective division, the point $(x, y, z, 1)$ goes to

$$x'' = x/z$$

$$y'' = y/z$$

$$Z'' = -(\alpha + \beta/z)$$

which projects orthogonally to the desired point
regardless of α and β



Picking α and β

If we pick

$$Z'' = -(\alpha + \beta/z)$$

$$\alpha = \frac{\text{near} + \text{far}}{\text{far} - \text{near}}$$

$$\beta = \frac{2\text{near} * \text{far}}{\text{near} - \text{far}}$$

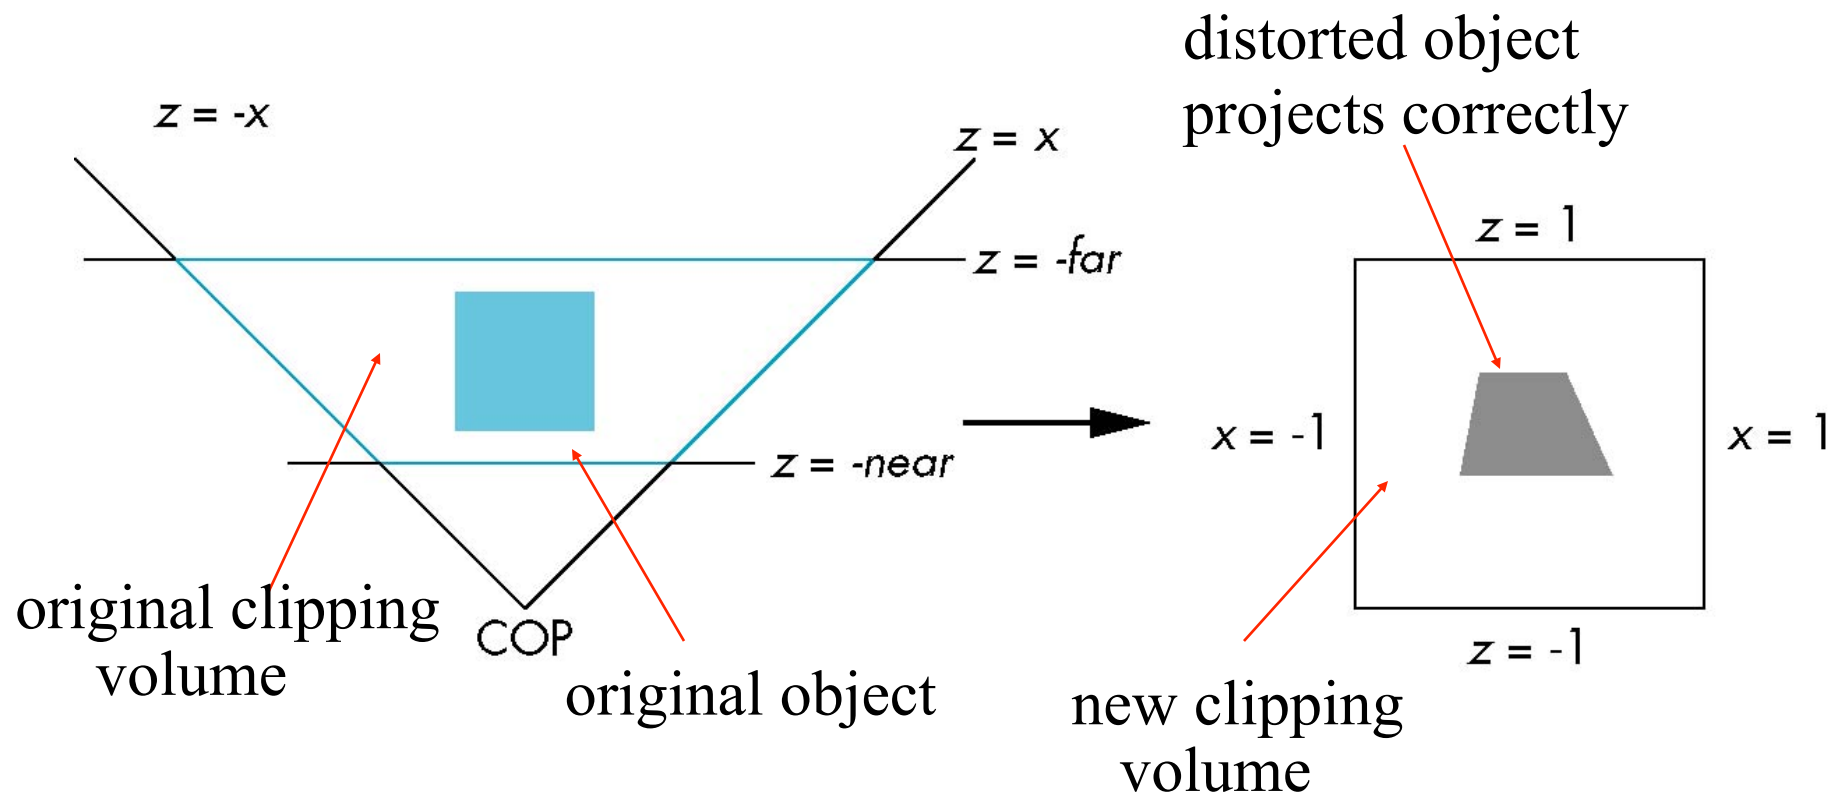
the near plane is mapped to $z = -1$

the far plane is mapped to $z = 1$

and the sides are mapped to $x = \pm 1, y = \pm 1$

Hence the new clipping volume is the default clipping volume

Normalization Transformation






Normalization and Hidden-Surface Removal

- Although our selection of the form of the perspective matrices may appear somewhat arbitrary, it was chosen so that if $z_1 > z_2$ in the original clipping volume then the for the transformed points $z_1' > z_2'$
- Thus hidden surface removal works if we first apply the normalization transformation
- However, the formula $z' = -(\alpha + \beta/z)$ implies that the distances are distorted by the normalization which can cause numerical problems especially if the near distance is small

OpenGL Perspective

- **Frustum** allows for an unsymmetric viewing frustum (although **Perspective** does not)

Frustum(left, right, bottom, top, near, far)




OpenGL Perspective

- **Perspective** provides less flexible, but more intuitive perspective viewing

`Perspective(fov, aspect, near, far);`

- Field of view in angles
- aspect: w/h
- $\text{near} < \text{far}$, both positive





OpenGL Perspective Matrix

- The normalization in **Frustum** requires an initial shear to form a right viewing pyramid, followed by a scaling to get the normalized perspective volume. Finally, the perspective matrix results in needing only a final orthographic transformation

$$\mathbf{P} = \mathbf{NSH}$$

our previously defined
perspective matrix

shear and scale



Why do we do it this way?

- Normalization allows for a single pipeline for both perspective and orthogonal viewing
- We stay in four dimensional homogeneous coordinates as long as possible to retain three-dimensional information needed for hidden-surface removal and shading
- We simplify clipping



Viewing OpenGL Code

- It's not too bad, thanks to Ed Angel
- In application

```
model_view = LookAt(eye, at, up);  
projection = Ortho(left, right, bottom, top,  
                  near, far);
```

or

```
projection = Perspective( fov, aspect, near,  
                          far);
```

- In vertex shader

```
gl_Position = projection*model_view*vPosition;
```



OpenGL code

- Remember that matrices are column major order in GLSL, so ...

Transpose your matrices when sending them to the shaders!



```
glUniformMatrix4fv(matrix_loc, 1, GL_TRUE,  
                    model_view);
```